

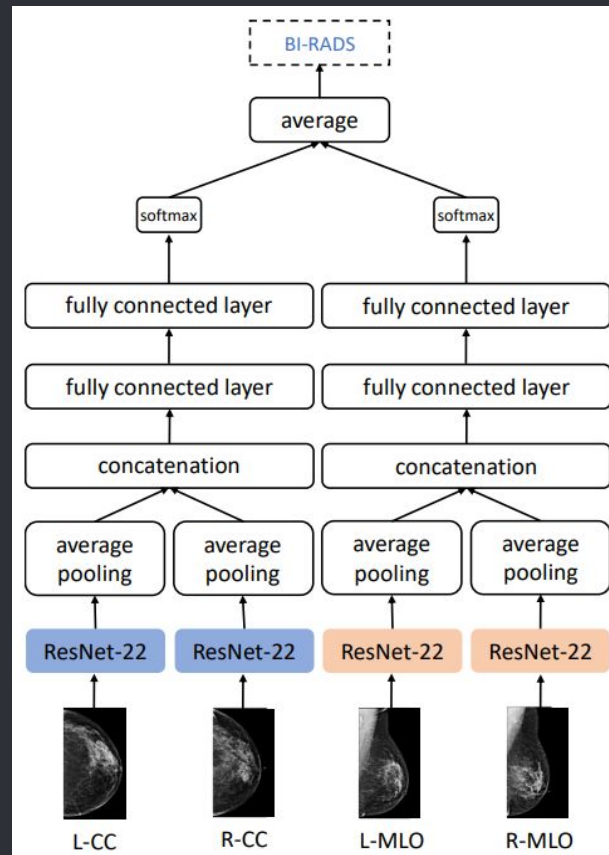
Medical Imaging RANZCR CLIP

Paolo Sofia, Matteo Carollo



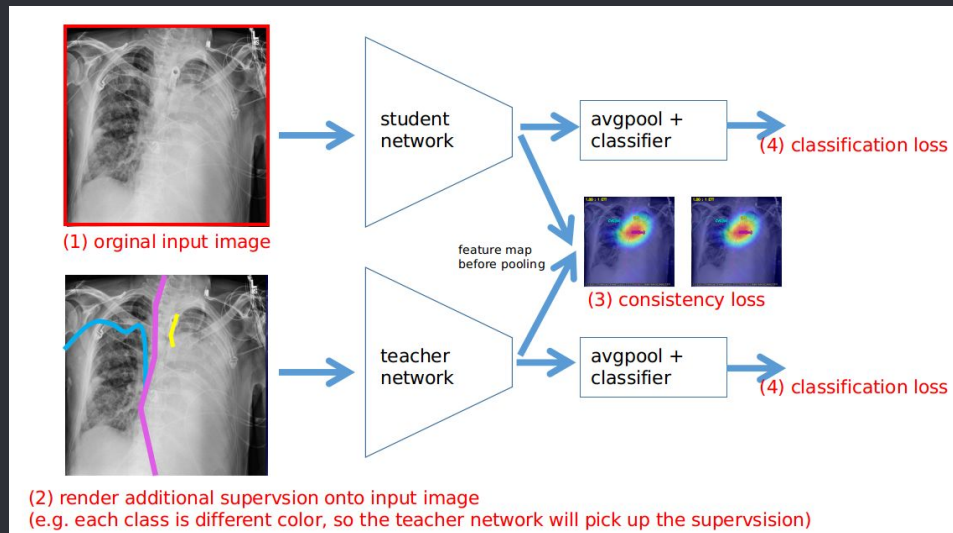
- Similar approaches - classification

We started our work taking into consideration this initial approach. Unfortunately, due to the limited time available, we simplified this approach, using a single model and attached a different head for each label group.



Similar approaches: distillation

We took inspiration also from other proposed method on kaggle, such as this one. We simplified our model by training a simple siamese net and then we fine tuned the model.



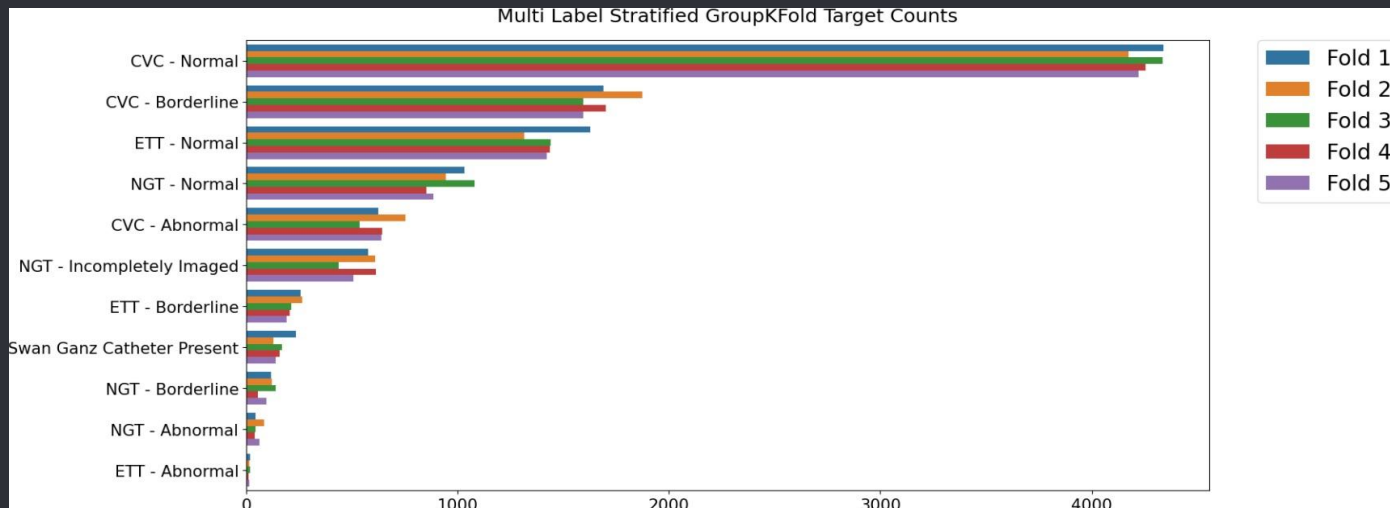
• Our work - Dataset

Dataset issue:

- Unbalanced labels
- Data leakage

Solution:

- Compute weights for each label
- Split train and validation set by patient maintaining the same label distribution



• Our work - Models, Losses, Optimizers comparison

Our initial approach was to compare different CNN, loss functions and optimizers to find the best combination of them

CNN

- Resnet18
- Resnet101
- Resnet50
- Inception V3
- Densenet169
- EfficientNetB2
- Xception

LOSS FUNCTIONS

- BCE
- Weighted BCE
- Weighted BCE + L1
- Weighted BCE + L2
- Weighted BCE + L1 L2
- ELR
- Focal Loss

OPTIMIZERS

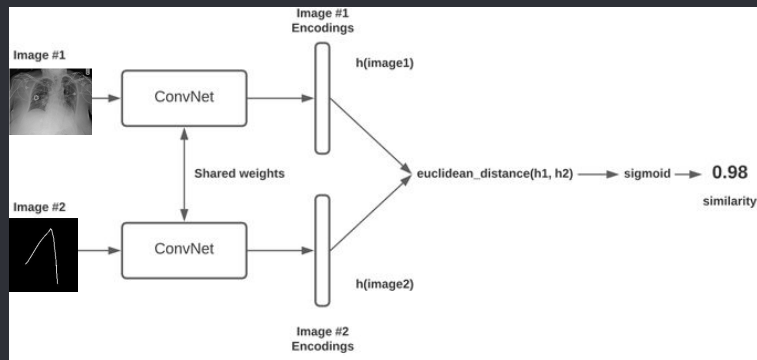
- Adam
- Adam w. Decaying Momentum
- SGD w. Decaying Momentum
- Yogi
- Adabound
- Adabelief

- Our work - Other approaches

Siamese Net

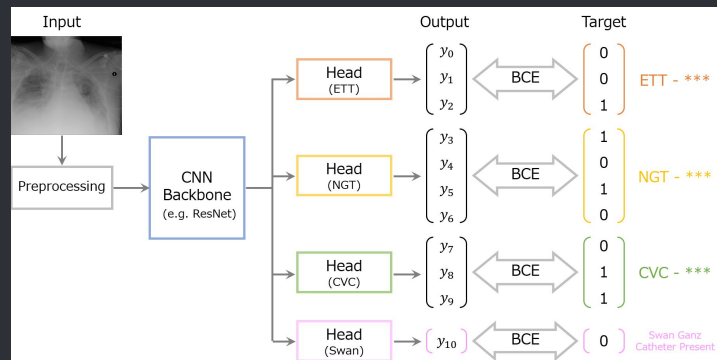
2 stages training

1. Train siamese net
2. Transfer learning siamese net weights to original net



Multi headed net

- Different Head for each group of labels

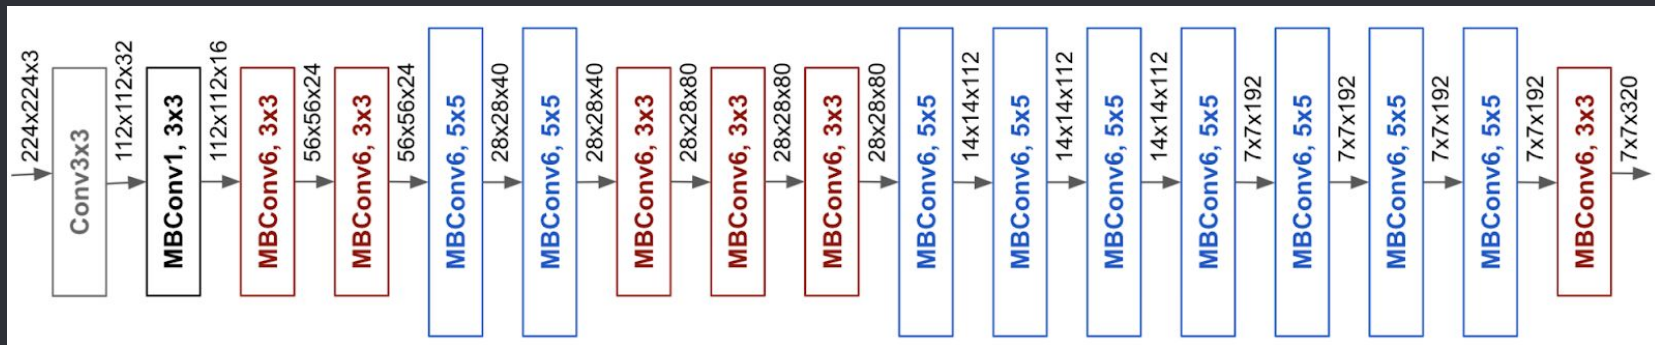
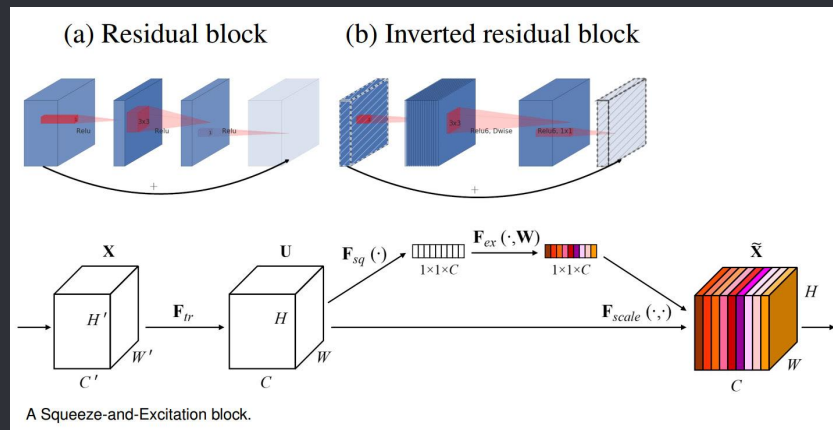


Architecture - EfficientNet B2

It has the best tradeoff between speed and performance

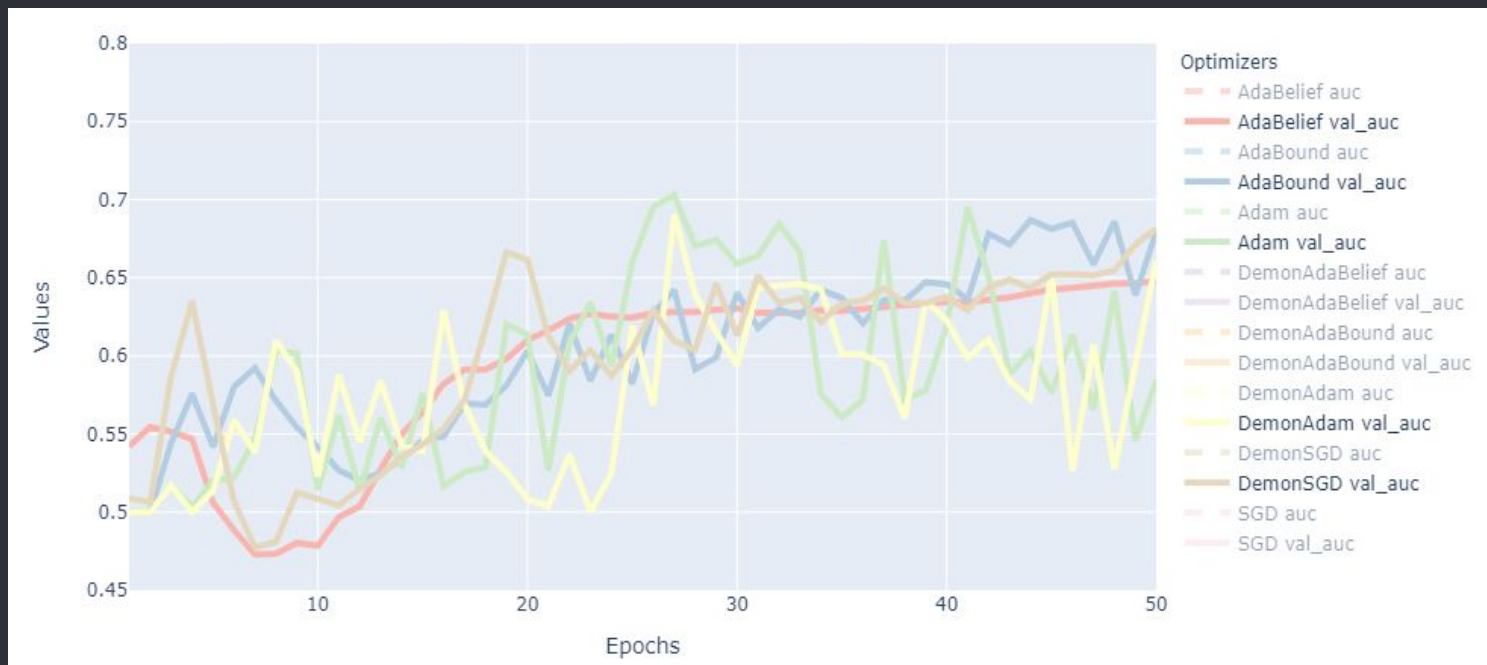
How it works:

- Inverted residual block
- Squeeze and excitation block



Results - Optimizer

- Adam is less stable, but it is faster to converge and achieves best results, so we chose it as optimizer. Then we used Adabelief for a second stage training as it behaves better in proximity of the minima



Results - Models

- EfficientNetB5 has slightly better performance than EfficientNetB2, but it is larger and slower, so we chose EfficientNetB2 also to avoid overfitting.



Results - EfficientNet B2

○ AUC


○ Training 0.9575

○ Validation 0.93

○ Area MCC

○ Training 0.5585

○ Validation 0.4962



● Conclusion

- The model used is very simple, but performance are good, even though, with better design of either siamese net or multi headed net, higher scores could be achieved.
- Compared to other models on kaggle, our performance may appear inferior, but on kaggle they focus on winning competition, ignoring three important aspects:
 - Data leakage
 - Loss functions
 - Noise as data augmentation technique

- Future work

Instance segmentation

Use the dataset annotations to use instance segmentation and then classify the segmented images

Distillation methods

Use teacher-student model using also annotation data

Multi headed net

Improve the design of the head, or use a different model for each group of data and then concatenate the results. In this case the models does not share the weights



Thanks for your attention

